



Introduction to Neural Networks

Data Boot Camp
Lesson 21.3



Google's DeepMind to Scan a Million Eyes to Fight Blindness with NHS

Google's DeepMind to Scan a Million Eyes to Fight Blindness with NHS

A Neural Network that merges the style of two images

deepdreamgenerator.com



Artificial Neural Networks

If this is what comes to mind when you think *neural*, you're on the right track:




The development of neural nets were inspired by the arrangement of neurons in the human brain.



The neuron's **axon** is what it uses to send signals to other neurons.




Dendrites are sensors that receive signals from other neurons.



Artificial Neural Networks

Neural Networks



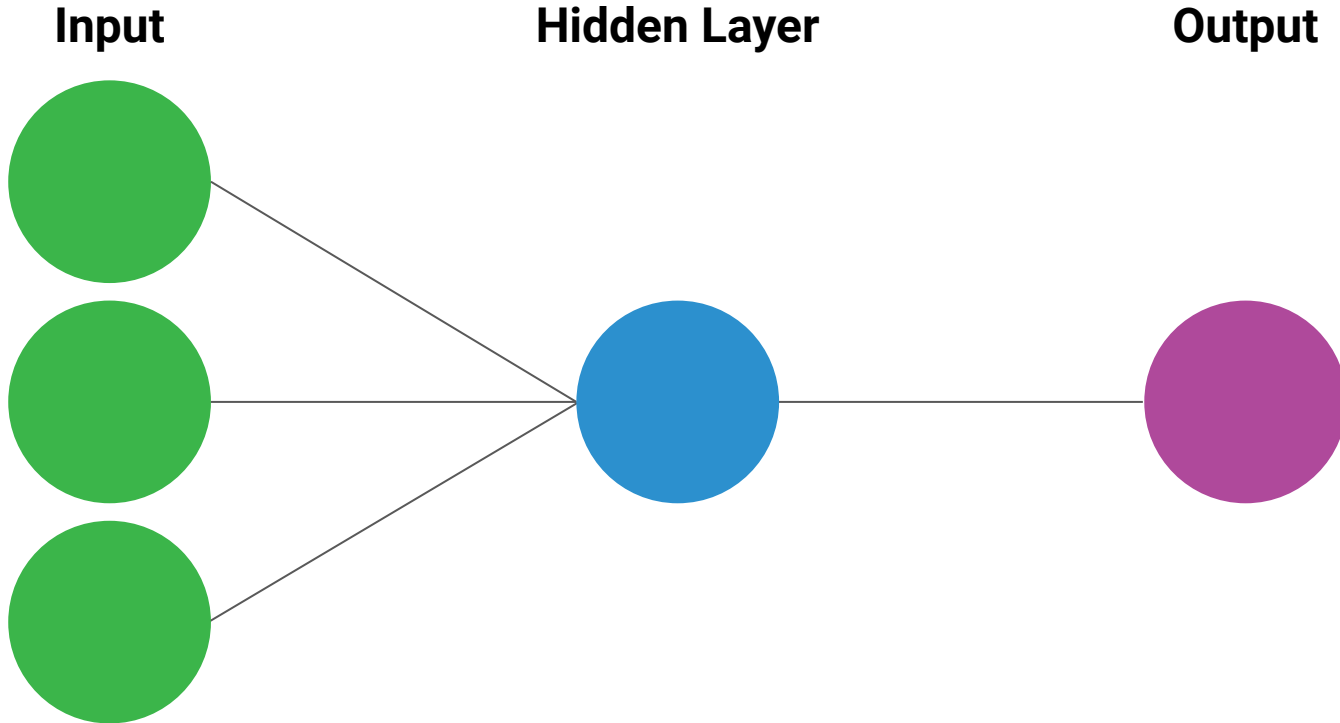
VS

The Brain



The Structure:

Neural networks are just layers of neurons connected together.




Artificial Neural Networks



As it turns out, the neural networks we program find patterns in inputs in much the same way the brain does.

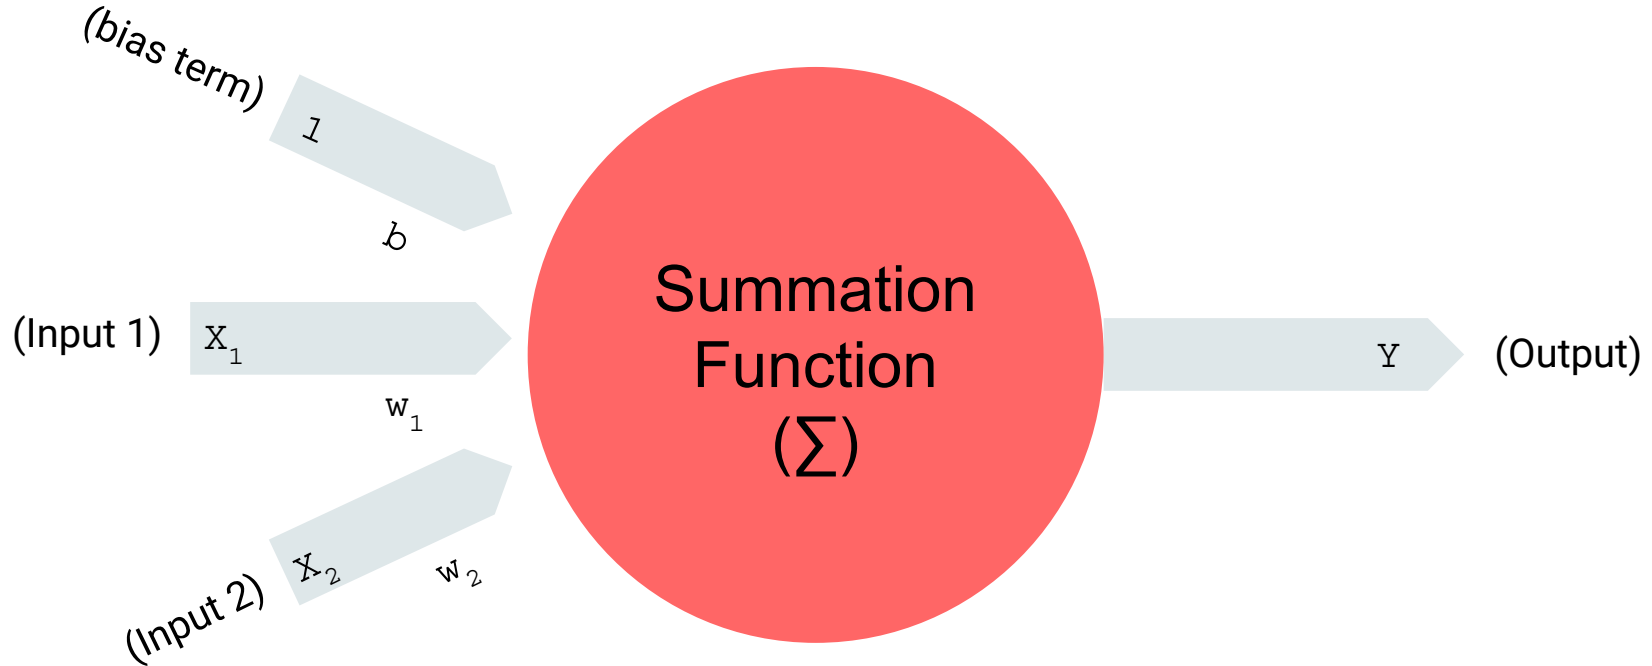


If you understand how the brain processes visual information, you'll understand how neural nets process *any* information.



The Math:

A single neuron transforms inputs into a mathematical output.



$$\text{Output of neuron} = Y = f(w_1x_1 + w_2x_2 + b)$$

The Math:




The most important mathematical detail to understand is that neural networks typically output **probabilities**.




In other words, if we were to tell it how many hours we studied for a test, and our grade on the midterm, it will spit out a *probability* that we'll pass (or fail!) the final.

Correct Output



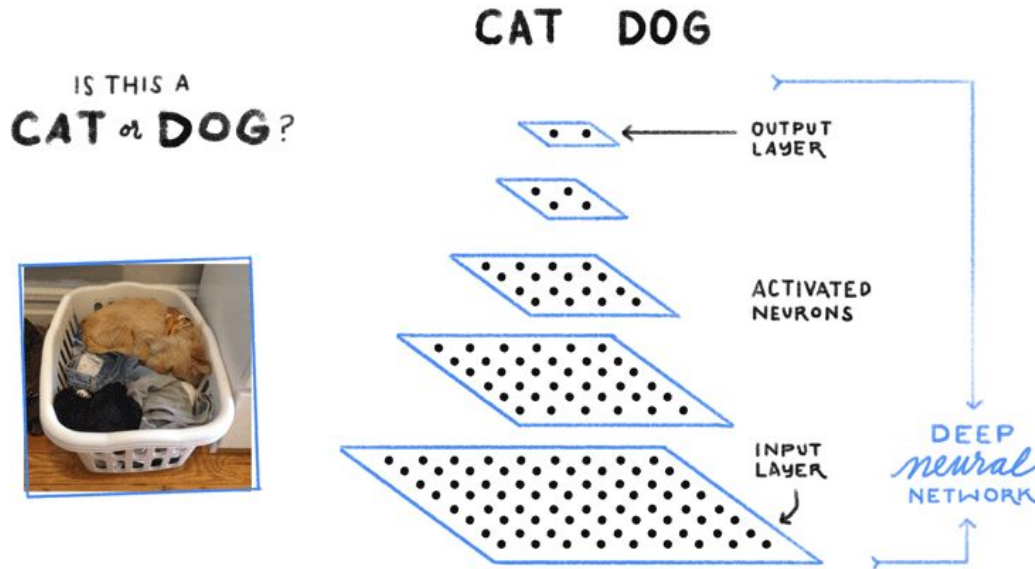
How Does It Work?

Neural networks are trained by feeding it (**forward propagation**) a bunch of inputs (like pictures of cats and dogs) and comparing the outputs to a known label (the output was a cat). The error in the output probability is then fed back (**back propagation**) to update the weights in each neuron. This is repeated until the neural network learns how to provide an accurate probability for a given input.



How Does It Work?

A trained neural network can take a new input (picture of a dog) and feed it through the network to predict the output label (cat or dog).





In a Nutshell:

If at first you don't
succeed, try, try again.



Questions?

Why Keras? What Is Keras?

Keras



A layer of abstraction



Simplifies deep learning modeling



Allows rapid prototyping



Works on top of TensorFlow (among other libraries)

TensorFlow



Machine learning library



Used to build neural networks



Runs on multiple platforms



Supports distributed computing



Allows detailed fine-tuning

Keras + TensorFlow

Keras allows interaction with TensorFlow through a simplified interface similar to Scikit-Learn.

Model → Fit → Predict (with a few other steps)

```
from tensorflow.keras.models import  
Sequential
```

```
model = Sequential()
```

```
# x_train and y_train are Numpy arrays -just like in the Scikit-Learn API.  
model.fit(x_train, y_train, epochs=5, batch_size=32)
```

```
classes = model.predict(x_test, batch_size=128)
```

Installation

Installation



conda
install
tensorflow

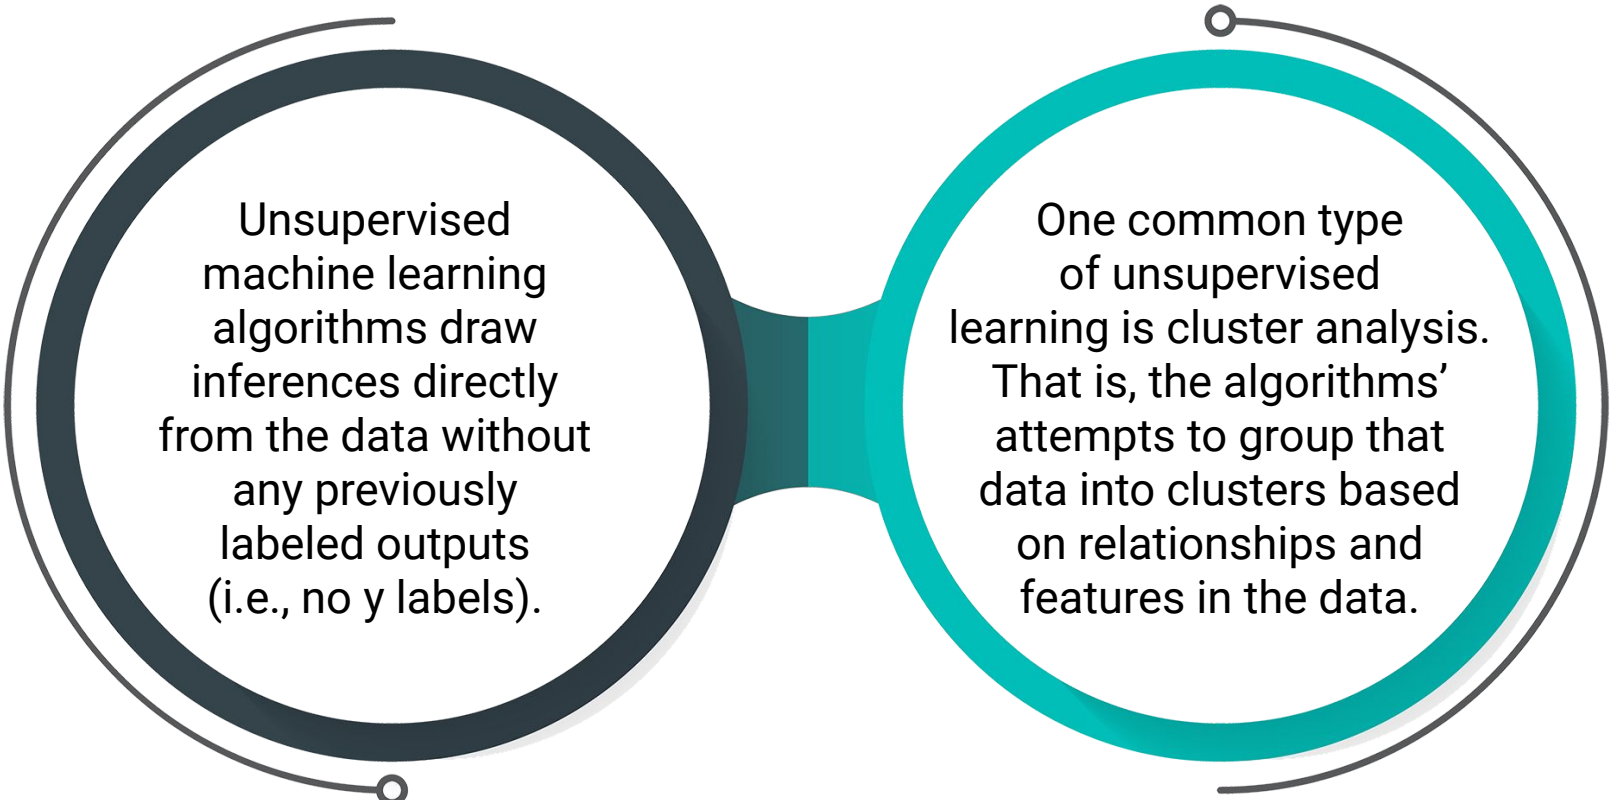
(or: pip
install
--upgrade
tensorflow)



Questions?

k-means

Unsupervised Learning



Unsupervised machine learning algorithms draw inferences directly from the data without any previously labeled outputs (i.e., no y labels).


The diagram consists of two large circles connected by a horizontal bar. The left circle has a dark blue border and contains text about unsupervised machine learning. The right circle has a teal border and contains text about cluster analysis. The connecting bar is split vertically, with the left half in dark blue and the right half in teal. Small circles at the top and bottom of the main circles' borders suggest a continuous path or flow.

One common type of unsupervised learning is cluster analysis. That is, the algorithms' attempts to group that data into clusters based on relationships and features in the data.


Clustering

While clustering may be intuitive to humans in this case, clustering algorithms have to decide which data points belong together.

Unclustered Data




Clustered Data



k-means Clustering Will Group the Data Into k Groups

The cluster center is the mean of all the points belonging to that cluster. Therefore, each point is closer to its own cluster's center than it is to other cluster centers.



k-means

k-means clustering will group the data into k number of groups.
A small k will create larger clusters.

K = 2



k-means

A large k will create smaller clusters.

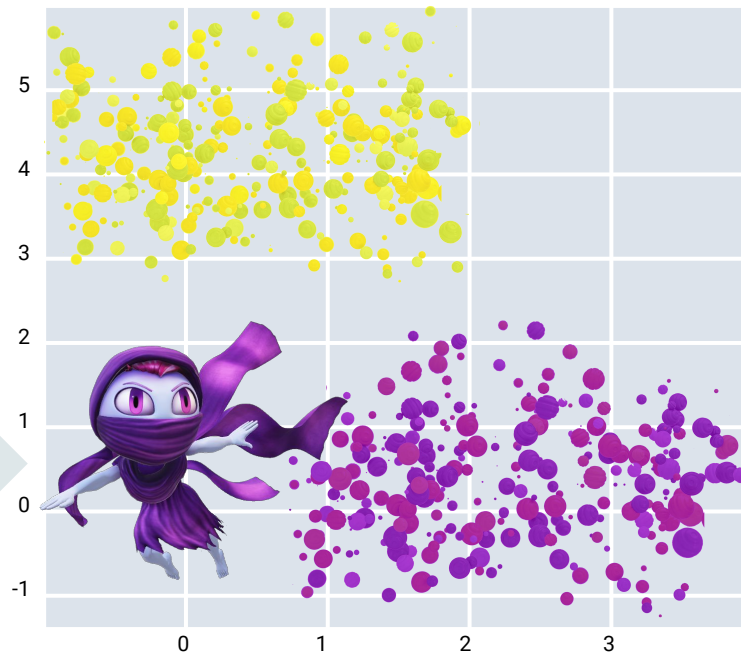
K = 5



k-means

Predictions using a trained k-means model

Clustered Data




Hi guys! Where do you want me?

Predictions Using a Trained k-means Model

What group would this new data point belong to?

Hi guys! Where do you want me?



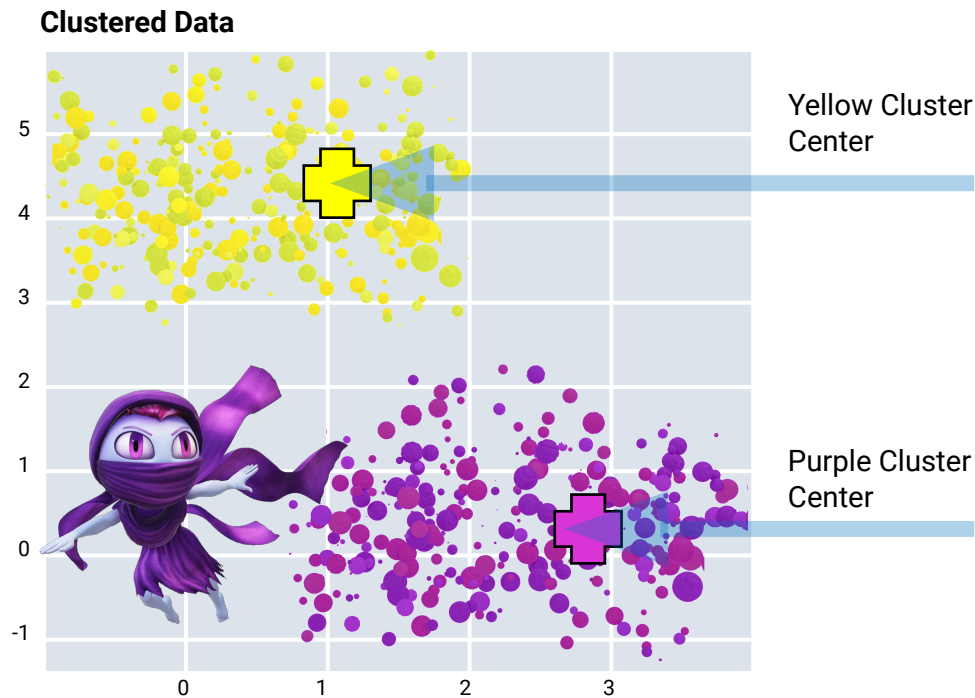
Predictions Using a Trained k-means Model

We should use Python to figure out which cluster center is the closest:

```
np.sqrt(sum((x - y) ** 2))
```


We can also figure this out using:

```
kmeans.predict(new_data)
```



Predictions Using a Trained k-means Model

Purple is closer!





Questions?